

Mathematical Background: Conditional Moment Conditions

1. Problem Formulation

1.1 Conditional Moment Restrictions

This library addresses estimation problems defined by **conditional moment restrictions (CMR)** of the form:

$$\mathbb{E}[\psi(X, Y; \theta) | Z] = 0 \quad P_Z\text{-almost surely}$$

where:

- (X, Y, Z) are random variables observed from an i.i.d. sample $\{(x_i, y_i, z_i)\}_{i=1}^n$
- $X \in \mathcal{X} \subseteq \mathbb{R}^{d_x}$ represents treatments or covariates
- $Y \in \mathcal{Y} \subseteq \mathbb{R}^{d_y}$ represents outcomes or responses
- $Z \in \mathcal{Z} \subseteq \mathbb{R}^{d_z}$ represents instrumental variables
- $\theta \in \Theta$ are the parameters of interest
- $\psi : \mathcal{X} \times \mathcal{Y} \times \Theta \rightarrow \mathbb{R}^k$ is a known moment function

1.2 Instrumental Variable Regression

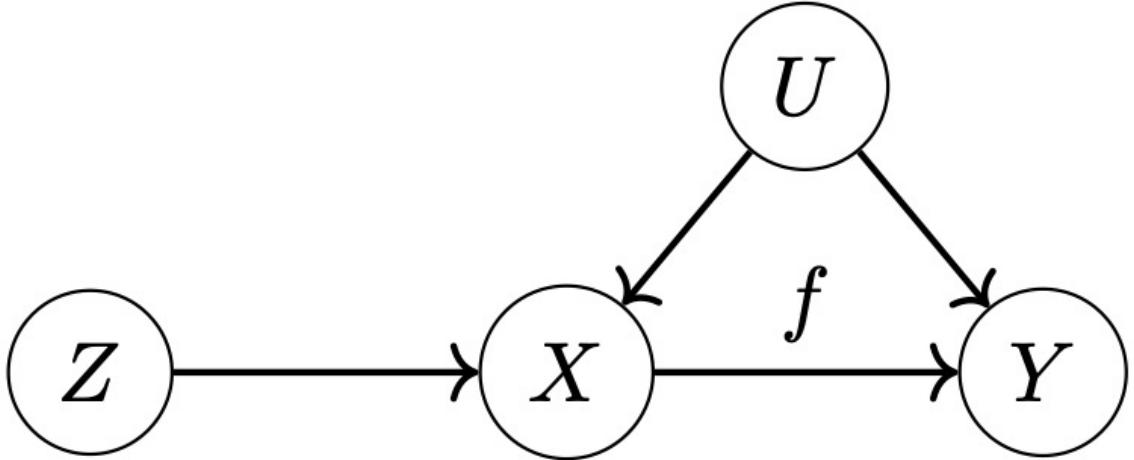
The canonical application is **instrumental variable (IV) regression**, where:

- We seek to estimate a structural function $f_\theta : \mathcal{X} \rightarrow \mathcal{Y}$
- The moment function takes the form: $\psi(X, Y; \theta) = f_\theta(X) - Y$
- The CMR becomes: $\mathbb{E}[f_\theta(X) - Y | Z] = 0$

This problem arises when:

- There exists unobserved confounding U between X and Y
- An instrument Z satisfies:
 1. **Relevance:** Z is correlated with X
 2. **Exclusion:** Z affects Y only through X
 3. **Exogeneity:** Z is independent of U

The causal structure is



1.3 Unconditional Moment Restrictions

When Z is not specified or the conditional restriction reduces to an unconditional one, we have the **moment restriction (MR)** problem:

$$\mathbb{E}[\psi(X, Y; \theta)] = 0$$

This is a special case where Z is trivial or integrated out.

2. Mathematical Framework

2.1 Dual Formulation via f-Divergences

The estimators in this library are best understood through the lens of **Generalized Empirical Likelihood (GEL)**, which formulates estimation as a constrained optimization problem on the space of probability measures.

2.1.1 The Primal Problem: Constrained Optimization Consider the unconditional moment restriction $\mathbb{E}[\psi(X, Y; \theta)] = 0$. The classical **Empirical Likelihood (EL)** estimator seeks a discrete probability distribution $p = (p_1, \dots, p_n)$ supported on the sample points that maximizes the likelihood (or equivalently minimizes the negative log-likelihood) subject to the moment constraints holding under p :

$$\begin{aligned} & \min_{p \in \mathbb{R}^n} \quad \sum_{i=1}^n -\log(np_i) \\ \text{subject to} \quad & \sum_{i=1}^n p_i \psi(x_i, y_i; \theta) = 0 \\ & \sum_{i=1}^n p_i = 1, \quad p_i \geq 0 \end{aligned}$$

The Parameter Explosion Problem: The primal problem attempts to estimate n parameters (the weights p_i) plus the parameters of interest θ . As $n \rightarrow \infty$, the number of nuisance parameters grows with the sample size, making direct optimization computationally infeasible and statistically challenging (“Neyman-Scott problem”).

2.1.2 The Dual Solution: Saddle Point Formulation To avoid this explosion, we use convex duality. This generalizes beyond EL to the class of **f-divergences**. For a convex function $f : \mathbb{R}_+ \rightarrow \mathbb{R}$ with $f(1) = 0$, we minimize the divergence between p and the uniform distribution $u_n = (1/n, \dots, 1/n)$:

$$\min_{p \in \Delta_n} \sum_{i=1}^n \frac{1}{n} f(np_i) \quad \text{s.t.} \quad \mathbb{E}_p[\psi(\theta)] = 0$$

By introducing Lagrange multipliers $\lambda \in \mathbb{R}^k$ for the moment constraints and using the Fenchel conjugate $f^*(t) = \sup_u \{tu - f(u)\}$, the problem converts to a dual formulation that depends only on λ and θ . This reduces the dimensionality from $O(n)$ to $O(k)$, leading to the **Generalized Empirical Likelihood (GEL)** saddle-point objective:

$$\min_{\theta \in \Theta} \sup_{\lambda \in \mathbb{R}^k} \left\{ -\frac{1}{n} \sum_{i=1}^n f^*(\lambda^\top \psi(x_i, y_i; \theta)) \right\}$$

This formulation is computationally efficient (k is fixed) and forms the basis for the modern machine learning approaches (VMM, FGEL, KMM) used in this library, which generalize λ from a vector to a function $h(Z)$.

3. Estimation Methods

3.1 Unconditional Moment Restrictions

3.1.1 Ordinary Least Squares (OLS) Objective:

$$\hat{\theta}_{\text{OLS}} = \arg \min_{\theta \in \Theta} \frac{1}{n} \sum_{i=1}^n \|\psi(x_i, y_i; \theta)\|^2$$

Use case: Simple baseline, assumes no confounding.

3.1.2 Generalized Method of Moments (GMM) Objective:

$$\hat{\theta}_{\text{GMM}} = \arg \min_{\theta \in \Theta} \bar{\psi}_n(\theta)^\top W_n \bar{\psi}_n(\theta)$$

where:

- $\bar{\psi}_n(\theta) = \frac{1}{n} \sum_{i=1}^n \psi(x_i, y_i; \theta)$
- W_n is a weighting matrix (typically $\hat{\Sigma}^{-1}$ where $\hat{\Sigma} = \frac{1}{n} \sum_{i=1}^n \psi_i \psi_i^\top$)

Two-step procedure:

1. Obtain initial estimate $\tilde{\theta}$
2. Compute W_n using $\tilde{\theta}$
3. Re-optimize with the optimal weighting matrix

Parameters: `reg_param` for regularization $W_n = (\hat{\Sigma} + \alpha I)^{-1}$

3.1.3 Generalized Empirical Likelihood (GEL) The library implements the dual formulation derived in Section 2.1.

Saddle-Point Objective:

$$\min_{\theta \in \Theta} \max_{\lambda \in \mathbb{R}^k} \mathcal{L}_{GEL}(\theta, \lambda) = \min_{\theta \in \Theta} \max_{\lambda \in \mathbb{R}^k} \left\{ -\frac{1}{n} \sum_{i=1}^n f^*(\lambda^\top \psi_i(\theta)) - \alpha \|\lambda\|^2 \right\}$$

Supported Divergences:

1. **Chi-squared (Euclidean EL):**

- Primal: Minimize $\sum(np_i - 1)^2$.
- Dual Conjugate: $f^*(t) = \frac{1}{2}(t+1)^2$.
- Connection: Closely related to Continuous Updating GMM (CUE).
- Code: `divergence='chi2'`

2. **Kullback-Leibler (Exponential Tilting):**

- Primal: Minimize $\sum p_i \log(np_i)$.
- Dual Conjugate: $f^*(t) = e^t$.
- Properties: Often more stable than EL; weights are strictly positive.
- Code: `divergence='kl'`

3. **Log-Euclidean (Empirical Likelihood):**

- Primal: Maximize $\sum \log(np_i)$.
- Dual Conjugate: $f^*(t) = -\log(1-t)$.
- Constraint: Requires $\lambda^\top \psi_i < 1$ for all i .
- Code: `divergence='log'`

Implementation: The `GeneralizedEL` class handles the optimization. The inner maximization over λ is solved via Newton's method (if `optim='lbfgs'`) or gradient ascent, while the outer minimization over θ uses the chosen optimizer (e.g., Adam, OAdam).

3.1.4 Kernel Method of Moments (KMM) **Formulation:** Minimizes an MMD-based objective with entropy regularization.

$$\min_{\theta} \max_{\nu, \eta, h} \left\{ \mathbb{E}_P[h(X, Y, Z)] + c - \gamma \mathbb{E}_Q[f^*(\frac{1}{\gamma}(h(X, Y, Z) + c - \nu^\top \psi(\theta)))] - \frac{\lambda}{2} \|h\|_{\mathcal{H}}^2 \right\}$$

where:

- P is the empirical distribution
- Q is a reference distribution (estimated via KDE)
- $h \in \mathcal{H}$ is an RKHS function on the joint space
- $\nu \in \mathbb{R}^k$ are dual variables for moments
- $c \in \mathbb{R}$ is a normalization constant

- $\gamma > 0$ is the entropy regularization parameter
- $\lambda > 0$ is the RKHS regularization parameter

Implementation details:

- Uses random Fourier features for h (dimension controlled by `n_random_features`)
- Reference distribution via KDE with bandwidth `kde_bandwidth`
- Samples from reference distribution: `n_reference_samples`
- Product kernel on (X, Y, Z) space

Parameters:

- `entropy_reg_param`: γ (main regularization)
- `rkhs_reg_param`: λ for RKHS norm
- `n_random_features`: Dimension D of RFF approximation
- `n_reference_samples`: Number of samples from reference distribution
- `kde_bandwidth`: Bandwidth for KDE

3.2 Conditional Moment Restrictions

The modern methods below generalize the GEL principle to **conditional** restrictions $\mathbb{E}[\psi | Z] = 0$. Instead of a single vector λ , we now learn a **function** of the instruments, effectively performing an “infinite” number of moment checks or learning the “optimal instrument” in a data-driven way.

3.2.1 Maximum Moment Restrictions (MMR) **Objective:** Directly minimizes the MMD between $\psi(\theta)$ and 0 conditional on Z .

$$\hat{\theta}_{\text{MMR}} = \arg \min_{\theta \in \Theta} \frac{1}{n^2} \sum_{i=1}^n \sum_{j=1}^n \psi(x_i, y_i; \theta)^\top K_{ij}^Z \psi(x_j, y_j; \theta)$$

where $K_{ij}^Z = k(z_i, z_j)$ is the kernel Gram matrix for instruments.

Advantages:

- Simple, no dual optimization
- Convex in θ for linear models
- Direct interpretation as MMD

Limitations:

- Can be sensitive to kernel choice
- No regularization on moment function

3.2.2 Sieve Minimum Distance (SMD) **Formulation:** Uses sieve basis (polynomials or splines) to approximate the conditional expectation.

Let $\{b_j(z)\}_{j=1}^J$ be a basis for $L^2(P_Z)$. Approximate:

$$\mathbb{E}[\psi(X, Y; \theta) | Z = z] \approx \sum_{j=1}^J \beta_j b_j(z)$$

Objective:

$$\min_{\theta, \beta} \frac{1}{n} \sum_{i=1}^n \left\| \psi(x_i, y_i; \theta) - \sum_{j=1}^J \beta_j b_j(z_i) \right\|^2$$

Properties:

- Classical semiparametric approach
- Theoretical guarantees under regularity conditions
- Basis dimension $J = J_n \rightarrow \infty$ as $n \rightarrow \infty$ (slower than n)

3.2.3 Variational Method of Moments (VMM) VMM generalizes GMM by replacing the fixed weighting matrix with a **learned instrument function** $h(Z)$ in an RKHS.

The Adversarial Game: Instead of minimizing a static quadratic form, VMM solves a minimax game where an “adversary” h tries to maximize the correlation with the moment violations, while the model θ tries to minimize it.

Kernel Version (VMM-kernel): The adversary is a function in a vector-valued RKHS \mathcal{H}^k . The objective simplifies analytically to:

$$\min_{\theta} \max_{h \in \mathcal{H}^k} \left\{ \frac{2}{n} \sum_{i=1}^n \psi(x_i, y_i; \theta)^{\top} h(z_i) - \|h\|_{\mathcal{H}^k}^2 \right\}$$

This leads to a closed-form solution for the optimal h , resulting in a generalized quadratic form involving the kernel matrix inverse:

$$\min_{\theta} \psi(\theta)^{\top} K^Z (K^Z K^Z + \alpha I)^{-1} K^Z \psi(\theta)$$

Neural Version (VMM-neural): Here, the adversary is a neural network $g_{\phi}(z)$. This recovers the **DeepGMM** method.

$$\min_{\theta} \max_{\phi} \left\{ \frac{2}{n} \sum_{i=1}^n \psi(\theta)^{\top} g_{\phi}(z_i) - \mathbb{E}[g_{\phi}(z)^{\top} g_{\phi}(z)] \right\}$$

3.2.4 Functional Generalized Empirical Likelihood (FGEL) FGEL is the direct conditional generalization of GEL. It replaces the vector λ in the GEL objective with a function $h(Z)$.

The Concept: We seek the “worst-case” re-weighting of the data (defined by $h(Z)$) that satisfies the conditional moment restrictions. The f-divergence acts as a regularizer on this re-weighting.

Objective:

$$\min_{\theta} \max_{h \in \mathcal{H}} \left\{ -\frac{1}{n} \sum_{i=1}^n f^*(h(z_i)^{\top} \psi(x_i, y_i; \theta)) - \frac{\alpha}{2} \|h\|_{\mathcal{H}}^2 \right\}$$

Kernel vs. Neural:

- **FGEL-kernel:** h lies in an RKHS. By the representer theorem, $h(z) = \sum \alpha_i k(z, z_i)$, reducing the problem to finding coefficients α_i .
- **FGEL-neural:** h is parametrized by a neural network. This allows for scalable estimation with large datasets where the kernel matrix would be too large to invert.

3.2.5 Kernel Method of Moments - Neural (KMM-neural) KMM-neural combines the Maximum Mean Discrepancy (MMD) principle with neural network test functions.

The Logic: It minimizes the MMD between the model distribution and a reference distribution, where the “test function” for the MMD is learned adversarially to detect moment violations.

Objective:

$$\min_{\theta} \max_{\nu, c, h} \left\{ \mathbb{E}_P[h] + c - \gamma \mathbb{E}_Q \left[f^* \left(\frac{1}{\gamma} (h + c - \nu_\phi(z)^\top \psi(\theta)) \right) \right] - \text{Reg}(h, \nu) \right\}$$

Key Innovation: Unlike FGEL which effectively reweights the *empirical* distribution, KMM introduces a reference distribution Q (e.g., from a Kernel Density Estimate). This stabilizes training and provides better finite-sample performance by ensuring the solution stays close to the data manifold. It uses **Random Fourier Features (RFF)** to scale the RKHS components to large datasets.

4. API Reference

4.1 High-Level Interface

estimation() Main entry point for training any estimator with automatic hyperparameter search.

```
from cmr import estimation

trained_model, stats = estimation(
    model,                      # torch.nn.Module
    train_data,                  # dict: {'t': ndarray, 'y': ndarray, 'z': ndarray}
    moment_function,             # callable: (model_output, y) -> torch.Tensor
    estimation_method,            # str: method name
    estimator_kwargs=None,       # dict: method-specific parameters
    hyperparams=None,             # dict: hyperparameter search space
    sweep_hparams=True,           # bool: whether to search hyperparameters
    validation_data=None,         # dict: validation set
    val_loss_func=None,            # callable or str: validation metric
    normalize_moment_function=True, # bool: normalize moments
    verbose=True                  # bool: print progress
)
```

Arguments:

Parameter	Type	Description
model	<code>torch.nn.Module</code>	Parametric model f_θ containing parameters to estimate
train_data	<code>dict</code>	Training data with keys ' <code>t</code> ' (treatments), ' <code>y</code> ' (outcomes), ' <code>z</code> ' (instruments)
moment_function	<code>callable</code>	Function $(f_\theta(x), y) \mapsto \psi(x, y; \theta) \in \mathbb{R}^k$
estimation_method	<code>str</code>	Method identifier (see table below)
estimator_kwargs	<code>dict</code>	Method-specific parameters (overrides defaults)
hyperparams	<code>dict</code>	Hyperparameter search space as <code>{param: [val1, val2, ...]}</code>
sweep_hparams	<code>bool</code>	If <code>True</code> , performs grid search over <code>hyperparams</code>
validation_data	<code>dict</code>	Validation set (if <code>None</code> , uses <code>train_data</code>)
val_loss_func	<code>callable or str</code>	Validation metric: custom function, ' <code>mmr</code> ', ' <code>hsic</code> ', or ' <code>moment_violation</code> '
normalize_moment_function	<code>bool</code>	If <code>True</code> , pretrains and normalizes moment components to unit variance
verbose	<code>bool or int</code>	Verbosity level: <code>False/0</code> , <code>True/1</code> , or <code>2</code>

Returns:

- `trained_model`: Trained PyTorch model (best from hyperparameter search)
- `stats`: Dictionary containing:
 - '`models
 - 'val_loss
 - 'hyperparam
 - 'best_index
 - 'train_stats`

Estimation Methods:

estimation_method	Problem Type	Description
'OLS'	Unconditional	Ordinary least squares
'GMM'	Unconditional	Generalized method of moments
'GEL'	Unconditional	Generalized empirical likelihood
'KMM'	Unconditional	Kernel method of moments
'SMD'	Conditional	Sieve minimum distance
'MMR'	Conditional	Maximum moment restrictions
'VMM-kernel'	Conditional	Variational MM with kernel
'VMM-neural'	Conditional	Variational MM with neural network

estimation_method	Problem Type	Description
'FGEL-kernel'	Conditional	Functional GEL with kernel
'FGEL-neural'	Conditional	Functional GEL with neural network
'KMM-neural'	Conditional	Kernel MM with neural instrument function

4.2 Low-Level Interface

For fine-grained control without hyperparameter search:

```
from cmr.methods.kmm_neural import KMMNeural

estimator = KMMNeural(
    model=model,
    moment_function=moment_function,
    entropy_reg_param=10.0,
    reg_param=1e-2,
    n_random_features=2000,
    verbose=True
)

estimator.train(train_data, validation_data)
trained_model = estimator.model
parameters = estimator.get_trained_parameters()
```

4.3 Data Format

All data must be provided as dictionaries with NumPy arrays:

```
train_data = {
    't': np.array([[...]]), # shape: (n_samples, d_t)
    'y': np.array([[...]]), # shape: (n_samples, d_y)
    'z': np.array([[...]]) # shape: (n_samples, d_z) or None for unconditional
}
```

For unconditional moment restrictions, set 'z': None.

4.4 Moment Function

The moment function must have signature:

```
def moment_function(model_output: torch.Tensor, y: torch.Tensor) -> torch.Tensor:
    """
    Args:
        model_output: Model predictions f_theta(x), shape (n, d_output)
        y: Observed outcomes, shape (n, d_y)
```

```

>Returns:
    Moment values  $\psi(x, y; \theta)$ , shape  $(n, k)$ 
"""
return model_output - y # Example: IV regression

```

Common moment functions:

1. **IV Regression**: lambda pred, y: pred - y
2. **Quantile Regression**: lambda pred, y: $(y - \text{pred}) * (\tau - (y < \text{pred}).\text{float}())$
3. **Vector-valued**: Return tensor with shape[1] = k for k moment conditions

4.5 Validation Metrics

Built-in validation metrics (string identifiers):

- 'moment_violation': $\frac{1}{n} \sum_{i=1}^n \|\psi(x_i, y_i; \theta)\|^2$ (unconditional)
- 'mmr': $\frac{1}{n^2} \sum_{i,j} \psi_i^\top K_{ij}^Z \psi_j$ (conditional, requires Z)
- 'hsic': Hilbert-Schmidt Independence Criterion between $\psi(\theta)$ and Z (conditional)

Custom validation functions:

```

def custom_val_loss(model: torch.nn.Module, val_data: dict) -> float:
"""
Args:
    model: Trained PyTorch model
    val_data: Validation data dict with keys 't', 'y', 'z'

>Returns:
    Scalar validation loss (lower is better)
"""
# Your custom validation logic
return loss_value

```

4.6 Common Workflow

```

import torch
import numpy as np
from cmr import estimation

# 1. Generate/load data
train_data = {'t': X_train, 'y': Y_train, 'z': Z_train}
val_data = {'t': X_val, 'y': Y_val, 'z': Z_val}
test_data = {'t': X_test, 'y': Y_test, 'z': Z_test}

# 2. Define model architecture
model = torch.nn.Sequential(
    torch.nn.Linear(dim_t, 50),
    torch.nn.LeakyReLU(),
    torch.nn.Linear(50, 20),
    torch.nn.LeakyReLU(),
    torch.nn.Linear(20, dim_y)

```

```

)
# 3. Define moment function
def moment_function(pred, y):
    return pred - y

# 4. Train with hyperparameter search
trained_model, stats = estimation(
    model=model,
    train_data=train_data,
    moment_function=moment_function,
    estimation_method='KMM-neural',
    validation_data=val_data,
    verbose=True
)

# 5. Evaluate
predictions = trained_model(torch.Tensor(test_data['t']))
mse = torch.mean((predictions - torch.Tensor(test_data['y']))**2)

```

5. Theoretical Guarantees

5.1 Identification

For the CMR to identify θ , we require:

- **Completeness:** The conditional expectation operator $\mathbb{E}[\cdot | Z]$ is injective
- **Global identification:** θ is the unique solution to $\mathbb{E}[\psi(X, Y; \theta) | Z] = 0$

Under completeness, the CMR is equivalent to:

$$\int \psi(x, y; \theta) p(x, y | z) dx dy = 0 \quad \text{for all } z \in \mathcal{Z}$$

5.2 Consistency and Convergence Rates

Under regularity conditions:

MMR:

- Consistency: $\hat{\theta}_n \xrightarrow{p} \theta_0$
- Rate: $\|\hat{\theta}_n - \theta_0\| = O_p(n^{-r})$ where r depends on smoothness and dimension

FGEL-kernel:

- With RKHS regularization $\alpha_n = O(n^{-\gamma})$ for appropriate γ
- Achieves minimax optimal rates under smoothness assumptions

KMM:

- Combines advantages of MMD testing and moment restriction estimation
- Reference distribution improves finite-sample performance
- Convergence rate depends on RFF approximation quality and reference sample size

5.3 Asymptotic Normality

For semiparametric efficient estimation, under regularity:

$$\sqrt{n}(\hat{\theta}_n - \theta_0) \xrightarrow{d} \mathcal{N}(0, V)$$

where V is the semiparametric efficiency bound. Methods like FGEL and VMM can achieve this bound with appropriate choices of divergence and regularization.

5.4 Computational Complexity

Method	Per-Iteration Complexity	Memory
OLS/GMM	$O(nk^2)$	$O(nk)$
GEL	$O(nk)$	$O(nk)$
MMR	$O(n^2k)$	$O(n^2)$
VMM-kernel	$O(n^3k)$ (matrix inversion)	$O(n^2k)$
FGEL-kernel	$O(n^2k^2)$	$O(nk)$
VMM-neural	$O(nk \cdot \text{NN})$	$O(nk)$
FGEL-neural	$O(nk \cdot \text{NN})$	$O(nk)$
KMM	$O(n^2)$ (full kernel)	$O(n^2)$
KMM-neural	$O(nD + n \cdot \text{NN})$	$O(nD)$

where NN denotes neural network forward/backward pass cost and D is the RFF dimension.

Scalability recommendations:

- $n < 1000$: Any method works well
- $1000 \leq n < 5000$: Use neural versions or KMM with RFF
- $n \geq 5000$: Mandatory to use neural networks or RFF approximations
- Use mini-batch training (`batch_size` parameter) for $n > 10000$

References

1. Kremer, H., Nemmour, Y., Schölkopf, B., & Zhu, J. J. (2023). Estimation Beyond Data Reweighting: Kernel Method of Moments. arXiv:2305.10898
2. Kremer, H., Zhu, J. J., Muandet, K., & Schölkopf, B. (2022). Functional Generalized Empirical Likelihood Estimation for Conditional Moment Restrictions. ICML 2022.
3. Bennett, A., Kallus, N., & Schnabel, T. (2020). Deep Generalized Method of Moments for Instrumental Variable Analysis. NeurIPS 2019.
4. Muandet, K., Fukumizu, K., Sriperumbudur, B., & Schölkopf, B. (2017). Kernel Mean Embedding of Distributions: A Review and Beyond. Foundations and Trends in Machine Learning.